

# Project Himanshu

January 8, 2019

## 1 Introduction

The goal of this kernel is to build a Keras cnn model to classify 12 kinds of plant seedlings.

We will use the V2 Plant Seedlings Dataset. This dataset contains 5,539 images distributed between 12 classes. The images show plant seedlings at different growth stages. Of the 12 classes, 3 classes are crop seedlings and 9 are weed seedlings.

The images are in different sizes. We will resize all images to 96x96 and use only 250 images from each class. We won't do any image augmentation.

This notebook will focus on:

Creating the folder structure that Keras generators need.

Creating generators to feed the images from the folders into the model.

Model building and training.

Assessing the quality of the model by generating a confusion matrix and a classification report

Results

This model will produce a validation accuracy that is greater than 90% and an F1 score of approximately 0.75.

```
In [ ]: from numpy.random import seed
        seed(101)
        from tensorflow import set_random_seed
        set_random_seed(101)

        import pandas as pd
        import numpy as np

        import tensorflow
        import keras

        from tensorflow.python.keras.models import Sequential
        from tensorflow.python.keras.layers import Dense, Dropout, Conv2D, MaxPooling2D
        , Flatten
        from tensorflow.python.keras.optimizers import Adam
        from tensorflow.python.keras.metrics import categorical_crossentropy
        from tensorflow.python.keras.preprocessing.image import ImageDataGenerator
        from tensorflow.python.keras.models import Model
```

```
from tensorflow.python.keras.callbacks import EarlyStopping, ReduceLROnPlateau
, ModelCheckpoint
```

```
import os
import cv2
```

```
import imageio
import skimage
import skimage.io
import skimage.transform
```

```
from sklearn.utils import shuffle
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
import itertools
import shutil
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # Number of samples we will have in each class.
        smpl_sz = 250
```

```
        # The images will all be resized to this size.
        img_sz = 96
```

```
In [3]: # Listing all the directories in the folder
        os.listdir('/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2')
```

```
Out[3]: ['Common wheat',
        'Scentless Mayweed',
        'Cleavers',
        'Shepherds Purse',
        'Loose Silky-bent',
        'Charlock',
        'Maize',
        'Black-grass',
        'Fat Hen',
        'Small-flowered Cranesbill',
        'Sugar beet',
        'Common Chickweed']
```

```
In [6]: # Create a new directory to store all available images
        all_images_dir = 'all_images_dir'
        os.mkdir(all_images_dir)
```

```
In [7]: # This code copies all images from their seperate folders into the same
        # folder called all_images_dir, changing the file name at the same time
        # as some of the files have similar names
```

```

folder_list = os.listdir
('/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2')

for folder in folder_list:

    # create a path to the folder
    path = '/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2/'
        + str(folder)

    # create a list of all files in the folder
    file_list = os.listdir(path)

    # move the 0 images to all_images_dir
    for fname in file_list:

        # source path to image
        src = os.path.join(path, fname)

        # Change the file name because many images have the same file name.
        # Add the folder name to the existing file name.
        new_fname = str(folder) + '_' + fname

        # destination path to image
        dst = os.path.join(all_images_dir, new_fname)
        # copy the image from the source to the destination
        shutil.copyfile(src, dst)

```

```

In [8]: # Get a list of all images in the all_images_dir folder.
image_list = os.listdir('all_images_dir')

# Creating a dataframe of all the images.
df_data = pd.DataFrame(image_list, columns=['image_id'])

df_data.head()

```

```

Out[8]:
      image_id
0  Common Chickweed_710.png
1  Scentless Mayweed_347.png
2  Scentless Mayweed_78.png
3  Scentless Mayweed_13.png
4      Maize_235.png

```

```

In [9]: # Each file name has this kind of format: E.g.
# Loose Silky-bent_377.png

# This function will extract the class name from the file name of each image.
def extract_target(x):
    # split into a list

```

```

a = x.split('_')
# the target is the first index in the list
target = a[0]

return target

# create a new column called 'target'
df_data['target'] = df_data['image_id'].apply(extract_target)

df_data.head()

```

Out[9]:

	image_id	target
0	Common Chickweed_710.png	Common Chickweed
1	Scentless Mayweed_347.png	Scentless Mayweed
2	Scentless Mayweed_78.png	Scentless Mayweed
3	Scentless Mayweed_13.png	Scentless Mayweed
4	Maize_235.png	Maize

In [10]: # Checking the class distribution

```

df_data['target'].value_counts()

```

Out[10]:

Loose Silky-bent	762
Common Chickweed	713
Scentless Mayweed	607
Small-flowered Cranesbill	576
Fat Hen	538
Sugar beet	463
Charlock	452
Cleavers	335
Black-grass	309
Shepherds Purse	274
Maize	257
Common wheat	253

Name: target, dtype: int64

In [11]: # Get a list of classes

```

target_list = os.listdir
('/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2')

for target in target_list:

    # Filter out a target and take a random sample
    df = df_data[df_data['target'] == target].sample(smpl_sz, random_state=101)

    # if it's the first item in the list
    if target == target_list[0]:
        df_sample = df

```

```

else:
    # Concat the dataframes
    df_sample = pd.concat([df_sample, df], axis=0).reset_index(drop=True)

```

```

In [13]: # Display the balanced classes.
df_sample['target'].value_counts()

```

```

Out[13]: Scentless Mayweed          250
Common Chickweed          250
Loose Silky-bent          250
Charlock                  250
Shepherds Purse           250
Fat Hen                   250
Cleavers                  250
Common wheat              250
Maize                     250
Sugar beet                250
Black-grass               250
Small-flowered Cranesbill 250
Name: target, dtype: int64

```

```

In [14]: # train_test_split

# stratify=y creates a balanced validation set.
y = df_sample['target']

df_train, df_val = train_test_split(df_sample, test_size=0.10, random_state=101
                                   , stratify=y)

print(df_train.shape)
print(df_val.shape)

```

```

(2700, 2)
(300, 2)

```

```

In [20]: print("\n Train set class distribution")
print(df_train['target'].value_counts())

print("\n Val set class distribution")
print(df_val['target'].value_counts())

```

```

Train set class distribution
Scentless Mayweed          225
Common Chickweed          225
Loose Silky-bent          225
Charlock                  225

```

Shepherds Purse	225
Fat Hen	225
Cleavers	225
Common wheat	225
Maize	225
Sugar beet	225
Black-grass	225
Small-flowered Cranesbill	225

Name: target, dtype: int64

Val set class distribution

Fat Hen	25
Loose Silky-bent	25
Common wheat	25
Cleavers	25
Common Chickweed	25
Shepherds Purse	25
Charlock	25
Sugar beet	25
Scentless Mayweed	25
Black-grass	25
Small-flowered Cranesbill	25
Maize	25

Name: target, dtype: int64

```
In [21]: #Creating a Direcotry Structure
         folder_list = os.listdir
         ('/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2')
```

folder\_list

```
Out[21]: ['Common wheat',
          'Scentless Mayweed',
          'Cleavers',
          'Shepherds Purse',
          'Loose Silky-bent',
          'Charlock',
          'Maize',
          'Black-grass',
          'Fat Hen',
          'Small-flowered Cranesbill',
          'Sugar beet',
          'Common Chickweed']
```

```
In [23]: #Checking the direcotry structure
         os.listdir('/home/himanshu/Desktop/MyFolder/Python/Keras/Project/NonsegmentedV2')
```

```
Out[23]: ['Common wheat',
          'Scentless Mayweed',
```

```
'Cleavers',  
'Shepherds Purse',  
'Loose Silky-bent',  
'Charlock',  
'Maize',  
'Black-grass',  
'Fat Hen',  
'Small-flowered Cranesbill',  
'Sugar beet',  
'Common Chickweed']
```

```
In [24]: # Create a new directory  
base_dir = 'base_dir'  
os.mkdir(base_dir)
```

```
#[CREATE FOLDERS INSIDE THE BASE DIRECTORY]
```

```
# now we create 2 folders inside 'base_dir':
```

```
# train_dir  
    # Maize  
    # Fat Hen  
    # Shepherds Purse  
    # Common Chickweed  
    # Cleavers  
    # Charlock  
    # Loose Silky-bent  
    # Small-flowered Cranesbill  
    # Black-grass  
    # Scentless Mayweed  
    # Sugar beet  
    # Common wheat
```

```
# val_dir  
    # Maize  
    # Fat Hen  
    # Shepherds Purse  
    # Common Chickweed  
    # Cleavers  
    # Charlock  
    # Loose Silky-bent  
    # Small-flowered Cranesbill  
    # Black-grass  
    # Scentless Mayweed  
    # Sugar beet  
    # Common wheat
```

```

# create a path to 'base_dir' to which we will join the names of the new folders
# train_dir
train_dir = os.path.join(base_dir, 'train_dir')
os.mkdir(train_dir)

# val_dir
val_dir = os.path.join(base_dir, 'val_dir')
os.mkdir(val_dir)

# [CREATE FOLDERS INSIDE THE TRAIN AND VALIDATION FOLDERS]

# create new folders inside train_dir

for folder in folder_list:

    folder = os.path.join(train_dir, str(folder))
    os.mkdir(folder)

# create new folders inside val_dir

for folder in folder_list:

    folder = os.path.join(val_dir, str(folder))
    os.mkdir(folder)

# check that the folders have been created

os.listdir('base_dir/train_dir')

```

```

Out[24]: ['Common wheat',
          'Scentless Mayweed',
          'Cleavers',
          'Shepherds Purse',
          'Loose Silky-bent',
          'Charlock',
          'Maize',
          'Black-grass',
          'Fat Hen',
          'Small-flowered Cranesbill',
          'Sugar beet',
          'Common Chickweed']

```

```

In [25]: # Set the id as the index in df_data
df_data.set_index('image_id', inplace=True)

```

```

In [26]: df_data.head()

```



```

Out[26]:
                                     target
image_id
Common Chickweed_710.png      Common Chickweed
Scentless Mayweed_347.png    Scentless Mayweed
Scentless Mayweed_78.png     Scentless Mayweed
Scentless Mayweed_13.png     Scentless Mayweed
Maize_235.png                 Maize

```

```

In [27]: # Getting a list of train and val images
train_list = list(df_train['image_id'])
val_list = list(df_val['image_id'])

# Transferring the train images

for image in train_list:

    # the id in the csv file does not have the .tif extension
    # therefore we add it here
    fname = image
    # get the label for a certain image
    folder = df_data.loc[image, 'target']

    # source path to image
    src = os.path.join(all_images_dir, fname)
    # destination path to image
    dst = os.path.join(train_dir, folder, fname)

    # resize the image and save it at the new location
    image = cv2.imread(src)
    image = cv2.resize(image, (img_sz, img_sz))
    # save the image at the destination
    cv2.imwrite(dst, image)

# Transfer the val images

for image in val_list:

    # the id in the csv file does not have the .tif extension
    # therefore we add it here
    fname = image
    # get the label for a certain image
    folder = df_data.loc[image, 'target']

    # source path to image

```

```

src = os.path.join(all_images_dir, fname)
# destination path to image
dst = os.path.join(val_dir, folder, fname)

# resize the image and save it at the new location
image = cv2.imread(src)
image = cv2.resize(image, (img_sz, img_sz))
# save the image at the destination
cv2.imwrite(dst, image)

```

## 2 Starting Modeling

```

In [28]: train_path = 'base_dir/train_dir'
        valid_path = 'base_dir/val_dir'

```

```

num_train_samples = len(df_train)
num_val_samples = len(df_val)
train_batch_size = 10
val_batch_size = 10

```

```

train_steps = np.ceil(num_train_samples / train_batch_size)
val_steps = np.ceil(num_val_samples / val_batch_size)

```

```

In [30]: datagen = ImageDataGenerator(rescale=1.0/255)

```

```

train_gen = datagen.flow_from_directory(train_path,
                                       target_size=(img_sz, img_sz),
                                       batch_size=train_batch_size,
                                       class_mode='categorical')

```

```

val_gen = datagen.flow_from_directory(valid_path,
                                     target_size=(img_sz, img_sz),
                                     batch_size=val_batch_size,
                                     class_mode='categorical')

```

```

# Note: shuffle=False causes the test dataset to not be shuffled
test_gen = datagen.flow_from_directory(valid_path,
                                       target_size=(img_sz, img_sz),
                                       batch_size=1,
                                       class_mode='categorical',
                                       shuffle=False)

```

```

Found 2700 images belonging to 12 classes.
Found 300 images belonging to 12 classes.
Found 300 images belonging to 12 classes.

```

### 3 Creating the model Architecture

In [31]: # Source: <https://www.kaggle.com/fmarazzi/baseline-keras-cnn-roc-fast-5min-0-8253-lb>

```
kernel_size = (3,3)
pool_size= (2,2)
first_filters = 32
second_filters = 64
third_filters = 128

dropout_conv = 0.3
dropout_dense = 0.3

model = Sequential()
model.add(Conv2D(first_filters, kernel_size, activation = 'relu',
                 input_shape = (img_sz, img_sz, 3)))
model.add(Conv2D(first_filters, kernel_size, activation = 'relu'))
model.add(Conv2D(first_filters, kernel_size, activation = 'relu'))
model.add(MaxPooling2D(pool_size = pool_size))
model.add(Dropout(dropout_conv))

model.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model.add(MaxPooling2D(pool_size = pool_size))
model.add(Dropout(dropout_conv))

model.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model.add(MaxPooling2D(pool_size = pool_size))
model.add(Dropout(dropout_conv))

model.add(Flatten())
model.add(Dense(256, activation = "relu"))
model.add(Dropout(dropout_dense))
model.add(Dense(12, activation = "softmax"))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 94, 94, 32)	896
conv2d_2 (Conv2D)	(None, 92, 92, 32)	9248

conv2d_3 (Conv2D)	(None, 90, 90, 32)	9248
-----		
max_pooling2d_1 (MaxPooling2D)	(None, 45, 45, 32)	0
-----		
dropout_1 (Dropout)	(None, 45, 45, 32)	0
-----		
conv2d_4 (Conv2D)	(None, 43, 43, 64)	18496
-----		
conv2d_5 (Conv2D)	(None, 41, 41, 64)	36928
-----		
conv2d_6 (Conv2D)	(None, 39, 39, 64)	36928
-----		
max_pooling2d_2 (MaxPooling2D)	(None, 19, 19, 64)	0
-----		
dropout_2 (Dropout)	(None, 19, 19, 64)	0
-----		
conv2d_7 (Conv2D)	(None, 17, 17, 128)	73856
-----		
conv2d_8 (Conv2D)	(None, 15, 15, 128)	147584
-----		
conv2d_9 (Conv2D)	(None, 13, 13, 128)	147584
-----		
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 128)	0
-----		
dropout_3 (Dropout)	(None, 6, 6, 128)	0
-----		
flatten_1 (Flatten)	(None, 4608)	0
-----		
dense_1 (Dense)	(None, 256)	1179904
-----		
dropout_4 (Dropout)	(None, 256)	0
-----		
dense_2 (Dense)	(None, 12)	3084
=====		
Total params: 1,663,756		
Trainable params: 1,663,756		
Non-trainable params: 0		
-----		

## 4 Training Model

```
In [32]: model.compile(Adam(lr=0.0001), loss='binary_crossentropy',
                      metrics=['accuracy'])
```

```
In [33]: filepath = "model.h5"
         checkpoint = ModelCheckpoint(filepath, monitor='val_acc', verbose=1,
```

```

save_best_only=True, mode='max')

reduce_lr = ReduceLROnPlateau(monitor='val_acc', factor=0.5, patience=3,
                               verbose=1, mode='max', min_lr=0.00001)

callbacks_list = [checkpoint, reduce_lr]

history = model.fit_generator(train_gen, steps_per_epoch=train_steps,
                              validation_data=val_gen,
                              validation_steps=val_steps,
                              epochs=20, verbose=1,
                              callbacks=callbacks_list)

Epoch 1/20
269/270 [=====>.] - ETA: 0s - loss: 0.2730 - acc: 0.9164
Epoch 00001: val_acc improved from -inf to 0.91694, saving model to model.h5
270/270 [=====] - 33s 123ms/step - loss: 0.2727 - acc: 0.9165 - val_loss: 0.2727
Epoch 2/20
269/270 [=====>.] - ETA: 0s - loss: 0.2256 - acc: 0.9181
Epoch 00002: val_acc improved from 0.91694 to 0.91861, saving model to model.h5
270/270 [=====] - 19s 72ms/step - loss: 0.2257 - acc: 0.9181 - val_loss: 0.2257
Epoch 3/20
269/270 [=====>.] - ETA: 0s - loss: 0.2025 - acc: 0.9218
Epoch 00003: val_acc did not improve
270/270 [=====] - 19s 70ms/step - loss: 0.2026 - acc: 0.9218 - val_loss: 0.2026
Epoch 4/20
269/270 [=====>.] - ETA: 0s - loss: 0.1909 - acc: 0.9248
Epoch 00004: val_acc improved from 0.91861 to 0.92056, saving model to model.h5
270/270 [=====] - 20s 72ms/step - loss: 0.1909 - acc: 0.9248 - val_loss: 0.1909
Epoch 5/20
269/270 [=====>.] - ETA: 0s - loss: 0.1786 - acc: 0.9293
Epoch 00005: val_acc improved from 0.92056 to 0.92972, saving model to model.h5
270/270 [=====] - 20s 73ms/step - loss: 0.1782 - acc: 0.9294 - val_loss: 0.1782
Epoch 6/20
269/270 [=====>.] - ETA: 0s - loss: 0.1656 - acc: 0.9324
Epoch 00006: val_acc improved from 0.92972 to 0.93278, saving model to model.h5
270/270 [=====] - 20s 75ms/step - loss: 0.1656 - acc: 0.9324 - val_loss: 0.1656
Epoch 7/20
269/270 [=====>.] - ETA: 0s - loss: 0.1541 - acc: 0.9371
Epoch 00007: val_acc improved from 0.93278 to 0.93389, saving model to model.h5
270/270 [=====] - 20s 72ms/step - loss: 0.1541 - acc: 0.9371 - val_loss: 0.1541
Epoch 8/20
269/270 [=====>.] - ETA: 0s - loss: 0.1437 - acc: 0.9423
Epoch 00008: val_acc improved from 0.93389 to 0.94028, saving model to model.h5
270/270 [=====] - 20s 74ms/step - loss: 0.1437 - acc: 0.9423 - val_loss: 0.1437
Epoch 9/20
269/270 [=====>.] - ETA: 0s - loss: 0.1346 - acc: 0.9459

```

Epoch 00009: val\_acc improved from 0.94028 to 0.94222, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.1345 - acc: 0.9460 - val\_loss: 0.1345  
Epoch 10/20  
269/270 [=====>.] - ETA: 0s - loss: 0.1258 - acc: 0.9479  
Epoch 00010: val\_acc improved from 0.94222 to 0.94528, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.1257 - acc: 0.9479 - val\_loss: 0.1257  
Epoch 11/20  
269/270 [=====>.] - ETA: 0s - loss: 0.1144 - acc: 0.9522  
Epoch 00011: val\_acc improved from 0.94528 to 0.94611, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.1144 - acc: 0.9521 - val\_loss: 0.1144  
Epoch 12/20  
269/270 [=====>.] - ETA: 0s - loss: 0.1071 - acc: 0.9546  
Epoch 00012: val\_acc did not improve  
270/270 [=====] - 20s 73ms/step - loss: 0.1071 - acc: 0.9546 - val\_loss: 0.1071  
Epoch 13/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0996 - acc: 0.9579  
Epoch 00013: val\_acc improved from 0.94611 to 0.94861, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.0995 - acc: 0.9579 - val\_loss: 0.0995  
Epoch 14/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0945 - acc: 0.9599  
Epoch 00014: val\_acc did not improve  
270/270 [=====] - 20s 73ms/step - loss: 0.0944 - acc: 0.9600 - val\_loss: 0.0944  
Epoch 15/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0865 - acc: 0.9648  
Epoch 00015: val\_acc improved from 0.94861 to 0.95056, saving model to model.h5  
270/270 [=====] - 20s 74ms/step - loss: 0.0863 - acc: 0.9649 - val\_loss: 0.0863  
Epoch 16/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0814 - acc: 0.9656  
Epoch 00016: val\_acc improved from 0.95056 to 0.95139, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.0816 - acc: 0.9655 - val\_loss: 0.0816  
Epoch 17/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0761 - acc: 0.9686  
Epoch 00017: val\_acc did not improve  
270/270 [=====] - 20s 75ms/step - loss: 0.0761 - acc: 0.9686 - val\_loss: 0.0761  
Epoch 18/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0714 - acc: 0.9699  
Epoch 00018: val\_acc did not improve  
270/270 [=====] - 20s 73ms/step - loss: 0.0713 - acc: 0.9699 - val\_loss: 0.0713  
Epoch 19/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0645 - acc: 0.9732  
Epoch 00019: val\_acc improved from 0.95139 to 0.95500, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.0644 - acc: 0.9732 - val\_loss: 0.0644  
Epoch 20/20  
269/270 [=====>.] - ETA: 0s - loss: 0.0602 - acc: 0.9745  
Epoch 00020: val\_acc improved from 0.95500 to 0.95556, saving model to model.h5  
270/270 [=====] - 20s 73ms/step - loss: 0.0601 - acc: 0.9746 - val\_loss: 0.0601

## 5 Evaluate Model using the Val Set

```
In [35]: model.metrics_names
```

```
Out[35]: ['loss', 'acc']
```

```
In [36]: # Print the validation loss and accuracy.
```

```
# Here the best epoch will be used.
model.load_weights('model.h5')

val_loss, val_acc = \
    model.evaluate_generator(test_gen,
                             steps=len(df_val))

print('val_loss:', val_loss)
print('val_acc:', val_acc)
```

```
val_loss: 0.11705364881115633
```

```
val_acc: 0.955555555621783
```

## 6 Plot the Training Curve

```
In [37]: # display the loss and accuracy curves
```

```
import matplotlib.pyplot as plt



acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.figure()

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
```

```
Out[37]: <Figure size 432x288 with 0 Axes>
```



<Figure size 432x288 with 0 Axes>



## 7 Make a prediction on the val set

We need these predictions to print the Confusion Matrix and calculate the F1 score.

```
In [38]: # make a prediction
         predictions = model.predict_generator(test_gen, steps=len(df_val), verbose=1)
```

```
300/300 [=====] - 1s 4ms/step
```

```
In [39]: predictions.shape
```

```
Out[39]: (300, 12)
```

```
In [40]: # This is how to check what index keras has internally assigned to each class.
         test_gen.class_indices
```

```
Out[40]: {'Black-grass': 0,
          'Charlock': 1,
          'Cleavers': 2,
          'Common Chickweed': 3,
          'Common wheat': 4,
          'Fat Hen': 5,
          'Loose Silky-bent': 6,
          'Maize': 7,
          'Scentless Mayweed': 8,
          'Shepherds Purse': 9,
          'Small-flowered Cranesbill': 10,
          'Sugar beet': 11}
```

```
In [41]: # Put the predictions into a dataframe.
         # The columns need to be ordered to match the output of the previous cell
```

```
class_dict = train_gen.class_indices
```

```
# Get a list of the dict keys.
```

```
cols = class_dict.keys()
```

```
df_preds = pd.DataFrame(predictions, columns=cols)
```

```
df_preds.head()
```

```
Out[41]:
```

	Black-grass	Charlock	Cleavers	Common Chickweed	Common wheat	Fat Hen	\
0	0.182492	0.000369	0.005700	0.000122	0.593382	0.013446	
1	0.468289	0.000008	0.000479	0.000002	0.160517	0.014424	
2	0.082417	0.000829	0.412537	0.000011	0.005066	0.303426	
3	0.000738	0.000164	0.002487	0.003655	0.000009	0.002779	
4	0.428237	0.000036	0.000265	0.000012	0.120407	0.000982	

	Loose Silky-bent	Maize	Scentless Mayweed	Shepherds Purse \
0	0.199016	0.001122	0.002284	0.000166
1	0.355947	0.000012	0.000113	0.000001
2	0.161298	0.000003	0.014975	0.000003
3	0.000249	0.000203	0.000005	0.000152
4	0.448889	0.000179	0.000566	0.000009

	Small-flowered Cranesbill	Sugar beet
0	0.000755	0.001146
1	0.000051	0.000157
2	0.011455	0.007981
3	0.857591	0.131968
4	0.000067	0.000349

## 8 Creating a Confusion Matrix

In [42]: *# Get the labels of the test images.*

```
test_labels = test_gen.classes
```

In [43]: *# Source: Scikit Learn website*  
*# http://scikit-learn.org/stable/auto\_examples/*  
*# model\_selection/plot\_confusion\_matrix.html#sphx-glr-auto-examples-model-*  
*# selection-plot-confusion-matrix-py*

```
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    # set the size of the figure here
    plt.figure(figsize=(15,10))

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
```

```

plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=80) # set x-axis text angle here
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()

```

```

In [44]: # argmax returns the index of the max value in a row
cm = confusion_matrix(test_labels, predictions.argmax(axis=1))

```

```

In [45]: # Define the labels of the class indices. These need to match the
# order shown above.
cm_plot_labels = cols

```

```

plot_confusion_matrix(cm, cm_plot_labels, title='Confusion Matrix')


```

Confusion matrix, without normalization

```

[[ 7  0  1  0  1  1 14  0  0  0  1  0]
 [ 0 23  0  0  0  0  0  1  1  0  0  0]
 [ 0  1 23  0  0  0  0  1  0  0  0  0]
 [ 0  0  0 19  1  0  0  2  1  2  0  0]
 [ 0  0  4  1 15  0  4  0  0  0  0  1]
 [ 0  0  1  0  1 17  3  0  0  0  1  2]
 [ 3  0  0  0  1  0 21  0  0  0  0  0]
 [ 0  0  2  0  1  2  0 14  3  0  1  2]
 [ 0  0  0  0  0  0  0  1 23  0  0  1]
 [ 0  1  0  1  0  1  0  0  1 17  2  2]
 [ 0  0  0  1  0  0  0  0  0  0 24  0]
 [ 1  0  0  0  2  3  0  2  2  0  1 14]]

```



```
In [46]: from sklearn.metrics import classification_report

# Generate a classification report

# Get the true labels
y_true = test_gen.classes

# For this to work we need y_pred as binary labels not as probabilities
y_pred_binary = predictions.argmax(axis=1)

report = classification_report(y_true, y_pred_binary, target_names=cm_plot_labels)

print(report)
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Black-grass	0.64	0.28	0.39	25
Charlock	0.92	0.92	0.92	25
Cleavers	0.74	0.92	0.82	25
Common Chickweed	0.86	0.76	0.81	25
Common wheat	0.68	0.60	0.64	25
Fat Hen	0.71	0.68	0.69	25
Loose Silky-bent	0.50	0.84	0.63	25
Maize	0.67	0.56	0.61	25
Scentless Mayweed	0.74	0.92	0.82	25
Shepherds Purse	0.89	0.68	0.77	25
Small-flowered Cranesbill	0.80	0.96	0.87	25
Sugar beet	0.64	0.56	0.60	25
micro avg	0.72	0.72	0.72	300
macro avg	0.73	0.72	0.71	300
weighted avg	0.73	0.72	0.71	300